

Tribhuvan University

Institute of Engineering

Pulchowk Campus



Lab report on

Shared Memory and Semaphores

Submitted by
Arpan Pokharel
075BCT015
Group: A

Submitted to
Department of Electronics and Computer Engineering

Submission Date: Feb 28,2022

Shared Memory and Semaphores

Theory

Shared memory or IPC is a mechanism in which two or more processes sharing a common segment of memory that they can both read to and write from to communicate with one another. It's just other part of process's address space just like calling malloc(). Because it's just memory, shared memory is the fastest IPC mechanism of them all.

Shared memory is well, a shared resource. Without somehow of letting the process that have access to it know if it's safe to read and write to the shared memory area. We're leaving our code open to race condition. This problem arises because these systems which use preemptive multitasking, where the OS takes care of managing when and for how long each process gets to execute. In such a situation, the contents of shared memory area can end up modified easily. One simple way of solving this problem is semaphores.

Semaphores

Semaphores are another IPC mechanism available when developing on UNIX. They allow different process to synchronise their access to certain resources.

The most common type is called a binary Semaphore because they have two states locked or unlocked.

When a process wants exclusive access to a resource, shared memory being an example

they attempt to lock the semaphore associated with that resource. If the semaphore they are attempting to lock is already locked, the caller is suspended, otherwise they are granted to lock. When we're finished using resource, we unlock the resource and any processes that have attempted to lock the semaphore in the meantime are woken up again to attempt the lock again. This way only one process can have access to the resource at once.

resource without anyone may attempting the process that have accessed it now it will be ready to use. If the system does not do this locking on a lock basis, then these systems will be able to access the resource at the same time. which uses preemptive multitasking, so here the operating system will decide which process gets to execute. In such long each process gets to execute. Through execution, the contents of shared memory can end up modified causing inconsistency. One simple way of solving this problem is semaphores.

semaphores are used in multiprogramming environments. Semaphores are used to coordinate the execution of multiple processes. They are used to coordinate the execution of multiple processes.

Two common types of semaphores are binary semaphores and general semaphores. Binary semaphores are used to coordinate the execution of two processes. They are used to coordinate the execution of two processes.

General semaphores are used to coordinate the execution of more than two processes. They are used to coordinate the execution of more than two processes.

Program 1:

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
#include<sys/shm.h>
const key_t SHMKEY = 7890;
const int PERMS = 0666;
char exFlag = 0;
Void main()
{
    int shmid;
    char* data;
    shmid = shmget(SHMKEY, 1, PERM | IPC_CREAT);
    If(shmid < 0)
        printf("can't get the shared memory");
    data = (char*)shmat(shmid, 0, 0);
    *data = 6;
    If(!fork())
    {
        printf("press any key to exit ... \n");
        scanf("l", &data, 1);
        printf("child exiting ... \n");
        exit(0);
    }
    else
    {
        while(*data)
            printf("Received l from child \n", data);
        shmdt(data);
        shmctl(shmid, IPC_RMID, 0);
        printf("parent exiting .. \n");
        exit(0);
    }
}
```

OUTPUT

Press any key to exit...
r
child exiting...
Received r from child
Parent exiting..

Discussion

The code above shows use of shared memory to exchange data in between child and parent process. First of all a shared memory segment is allocated using `shmget` which could be used by all interested process. And then thus created segment is mapped onto address 0. And then data is written through child process which is read by parent process which later detaches the segment which is later destroyed.

Program 2:

```
#include<stdio.h>
#include<sys/types.h>
#include<sys/ipc.h>
#include<sys/sem.h>
void main()
{
    int semid, key, nsem;
    key = (key - t) 0 x 20;
    nsem = 0;
    semid = 0;
    semid = semget(key, nsem, IPC_CREAT | 0666);
    printf("created semaphore with id: %d, semid: %d\n");
}
```

OUTPUT

created semaphore with id: -1

DISCUSSION

`semget()` system call gets a system semaphore set identifies, i.e. returns the system semaphore set identifies associated with the argument key. The system call creates a new set of '`nsem`' number of semaphores. Thus in

above lab code, no semaphore is created and -1 is returned by semget() which would have returned semaphore set identifier (a non-negative integer)

changing nsem = 1

OUTPUT

created semaphore with id: 0
----- Semaphore Arrays -----
Key sem id owner perms nsems
0x00000020 0 Arpan 666 1

Discussion

changing value of nsem from 0 to 1, one semaphore is created with identifier

Program 3

```
#include <sys/types.h>
#include <sys/ipc.h>
main()
{
    int semid, key, nsem, flag;
    nsem = 1;
    flag = 0666 | IPC_CREAT;
    for (i = 0; i < nsem; i++)
    {
        key = (key + i) ^ i;
        semid = semget(key, nsem, flag);
        if (semid > 0)
            printf("created semaphore with Id %d", semid);
        else
            printf("maximum number of semaphore set are %d (%d)", nsem, i);
    }
    exit(0);
}
```

Output

Created semaphore with ID =: 32769

Created semaphore with ID =: 65538

Created semaphore with ID =: 1048608

maximum number of semaphore set are 32

Discussion

- - - - - semaphore Arrays - - - - -

key	sem id	owner	perms	nsems
0x00000020	0	Arpan	666	1
0x00000010	32769	Arpan	666	1
;				
0x0000001F	1048608	Arpan	666	1

The maximum number of semaphore that can be created is 32. i.e. from 0x00000000 to 0x0000001F. Because a semaphore with key 0x00000020 is already created in previous program and so it returns some negative value as it cannot create one and loop PS terminated.

Program 4

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
```

```
const key_t SEMKEY = 4567;
```

```
const int perms = 0666;
```

```
char ex_flag = 0
```

```
void main()
```

```
{ int semid;
```

```
semid = semget(SEMKEY, 1, perms | IPC_CREAT);
```

```

if (semid < 0) {
    printf("can't create semaphore"); exit(1)
} if (!fork()) {
    char data;
    sembuf *sbuf = (sembuf *) malloc(sizeof(sembuf));
    sbuf->sem_num = 0;
    sbuf->sem_op = 1;
    sbuf->sem_flg = 0;
    printf("press any key to exit... \n");
    scanf("%c", &data);
    semop(semid, sbuf, 1);
    printf("child exiting");
    exit(0);
}
else {
    sembuf *sbuf = (sembuf *) malloc(sizeof(sem - buf));
    sbuf->sem_num = 0;
    sbuf->sem_op = -1;
    sbuf->sem_flg = 0;
    semop(semid, sbuf, 1);
    printf("child released semaphore \n");
    printf("parent exiting");
    exit(0);
}

```

Output

press any key to exit ..

a

child exiting ..

child released semaphore

parent exiting ..

Discussion

Here semaphore is released by the child, then the parent continues. Until then the parent will be waiting at the statement

```
sem-op( semid, S bud, 1);
```

Here sem-op is # in the child is +1 which means to increment semvalue by 1 and sem-op value in parent is -1 which means to decrease the value by 1.

If sem-op is less than 0 and sem_flag is not equal to IPC_NOWAIT and if the semaphore value is too low, that particular operation is blocked until another operation release the resource or modifies the value of semaphore such that the requested operation can proceed. In above program, the sem-op() in parent is blocked and is executed after the child increase the value of semaphore.

Conclusion

Thus, shared memory (shm) and semaphores (sem) can be used as a means of inter process communication.